

CPSC 5042: Week 2


Understanding Main Memory

**“640k ought to be enough for
anybody.”**


What is Main Memory?

- Main Memory is the memory unit that works directly with CPU, and in which data and instructions must reside.
- Also called RAM (Random Access Memory)
- Famous Computer Scientist Jay Forrester developed it in 1947 as a part of a project in MIT called Whirlwind. At the time, they called it “coincident current magnetic core storage”. It later was just called RAM


Logical vs Physical Memory




- A computer system has a single physical memory consisting of n words addressed by fixed-size physical addresses in the range $[0 : n-1]$



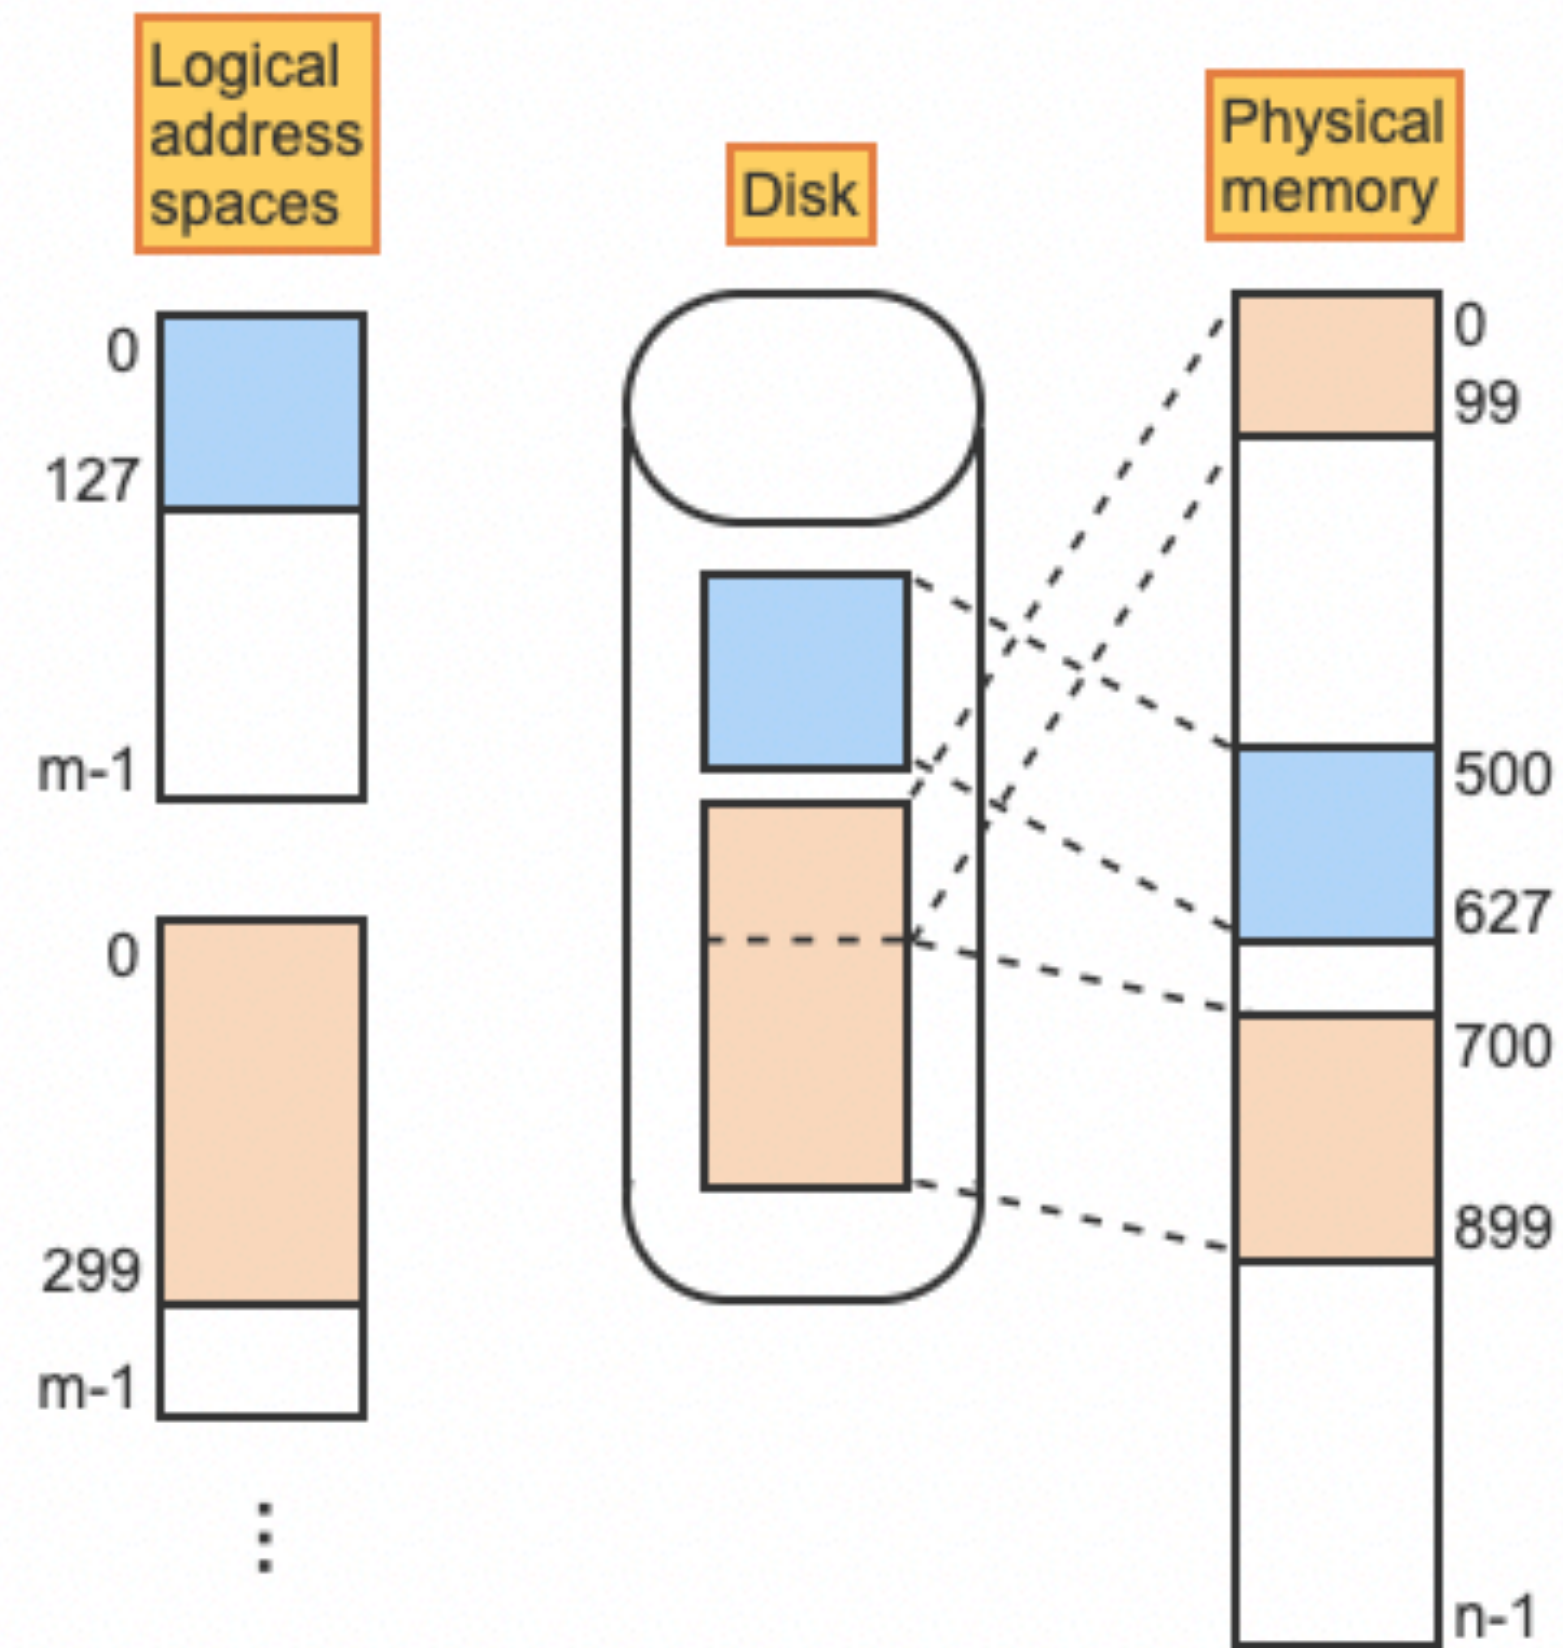
- To facilitate program development and program loading, multiple logical address spaces are provided with possibly different address ranges $[0 : m-1]$



- The occupied portion of each logical address space is kept on disk

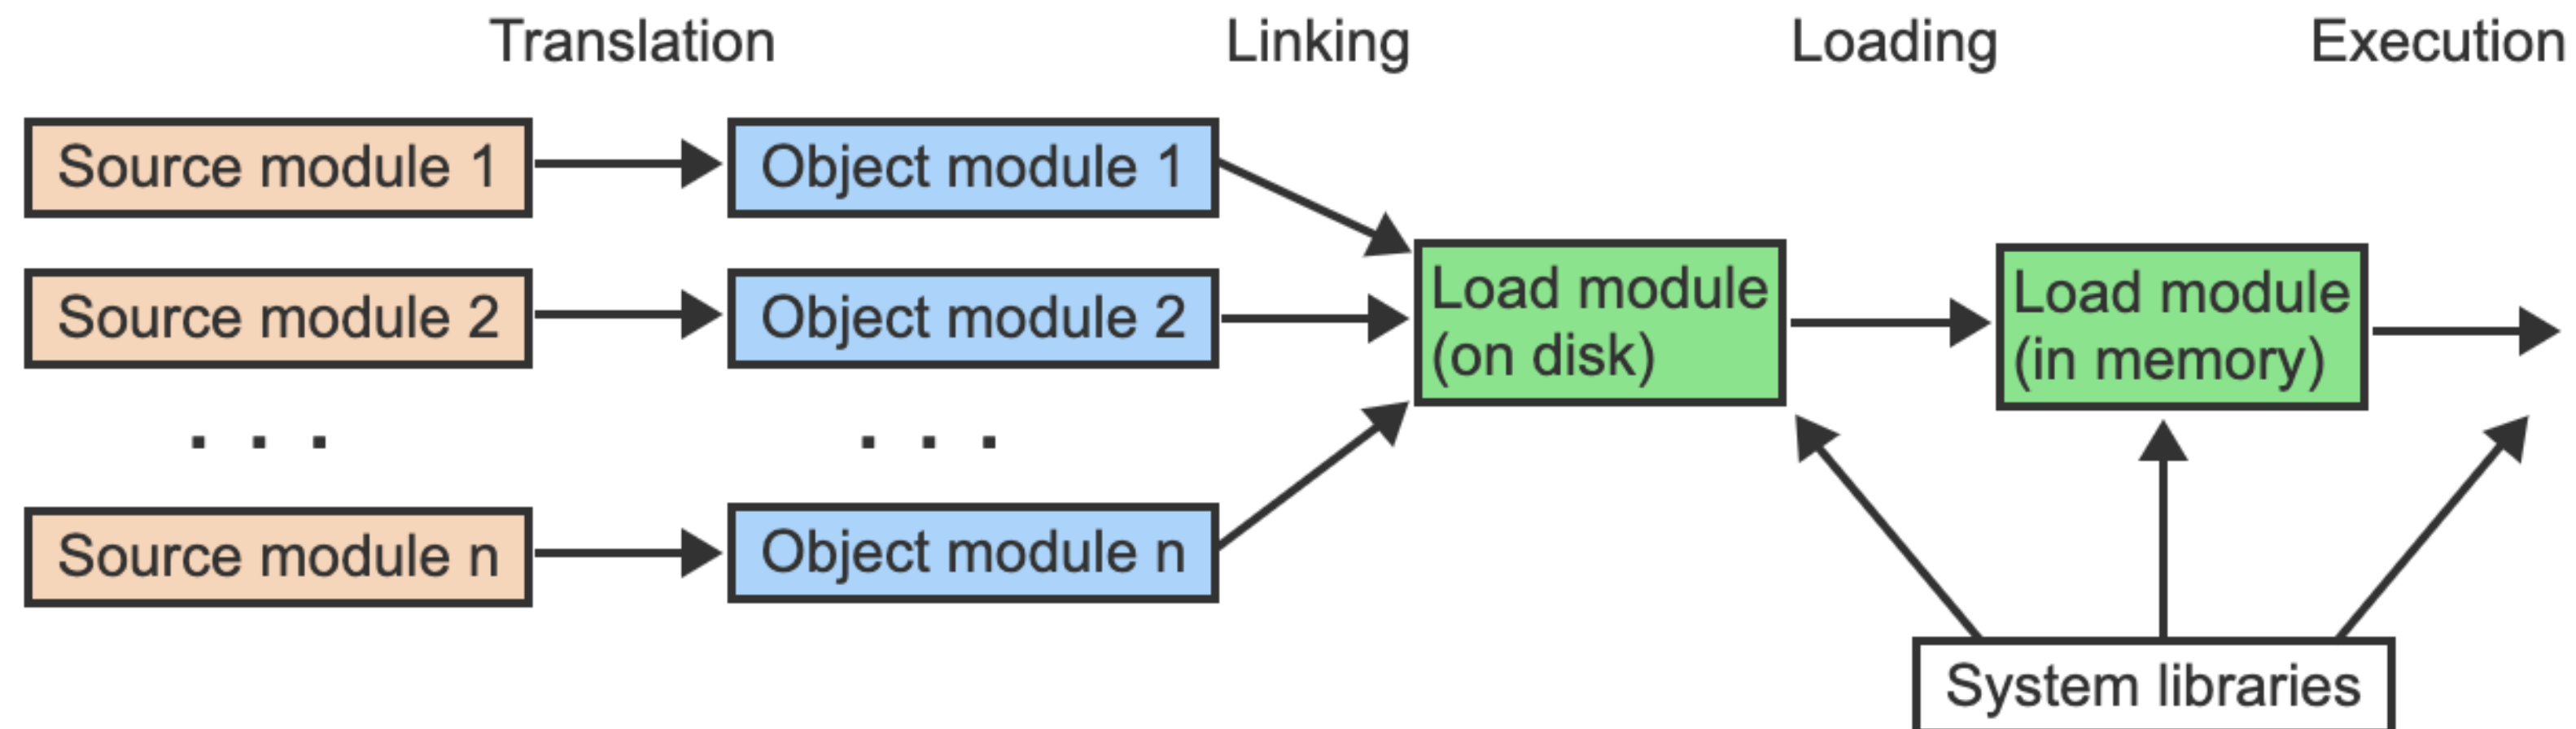


- A program may be copied to a contiguous region of physical memory when needed for execution. Logical space [0 : 127] is copied to memory locations [500 : 627]



- A program can also be copied to multiple disjoint regions of physical memory. The space [0 : 299] is mapped to the regions [0 : 99] and [700 : 899]

Program Transformations



- A program consists of one or more source modules written in a high-level language or assembly language.
- Each module is translated separately into an object module using a compiler or assembler.
- Object modules are combined into a load module by a linker. The load module is kept on disk.
- Prior to execution, the load module is copied to physical memory by a loader.
- Various system libraries, in the form of object modules, may be included in the load module by the linker, the loader, or during execution

Relocations & Address Binding

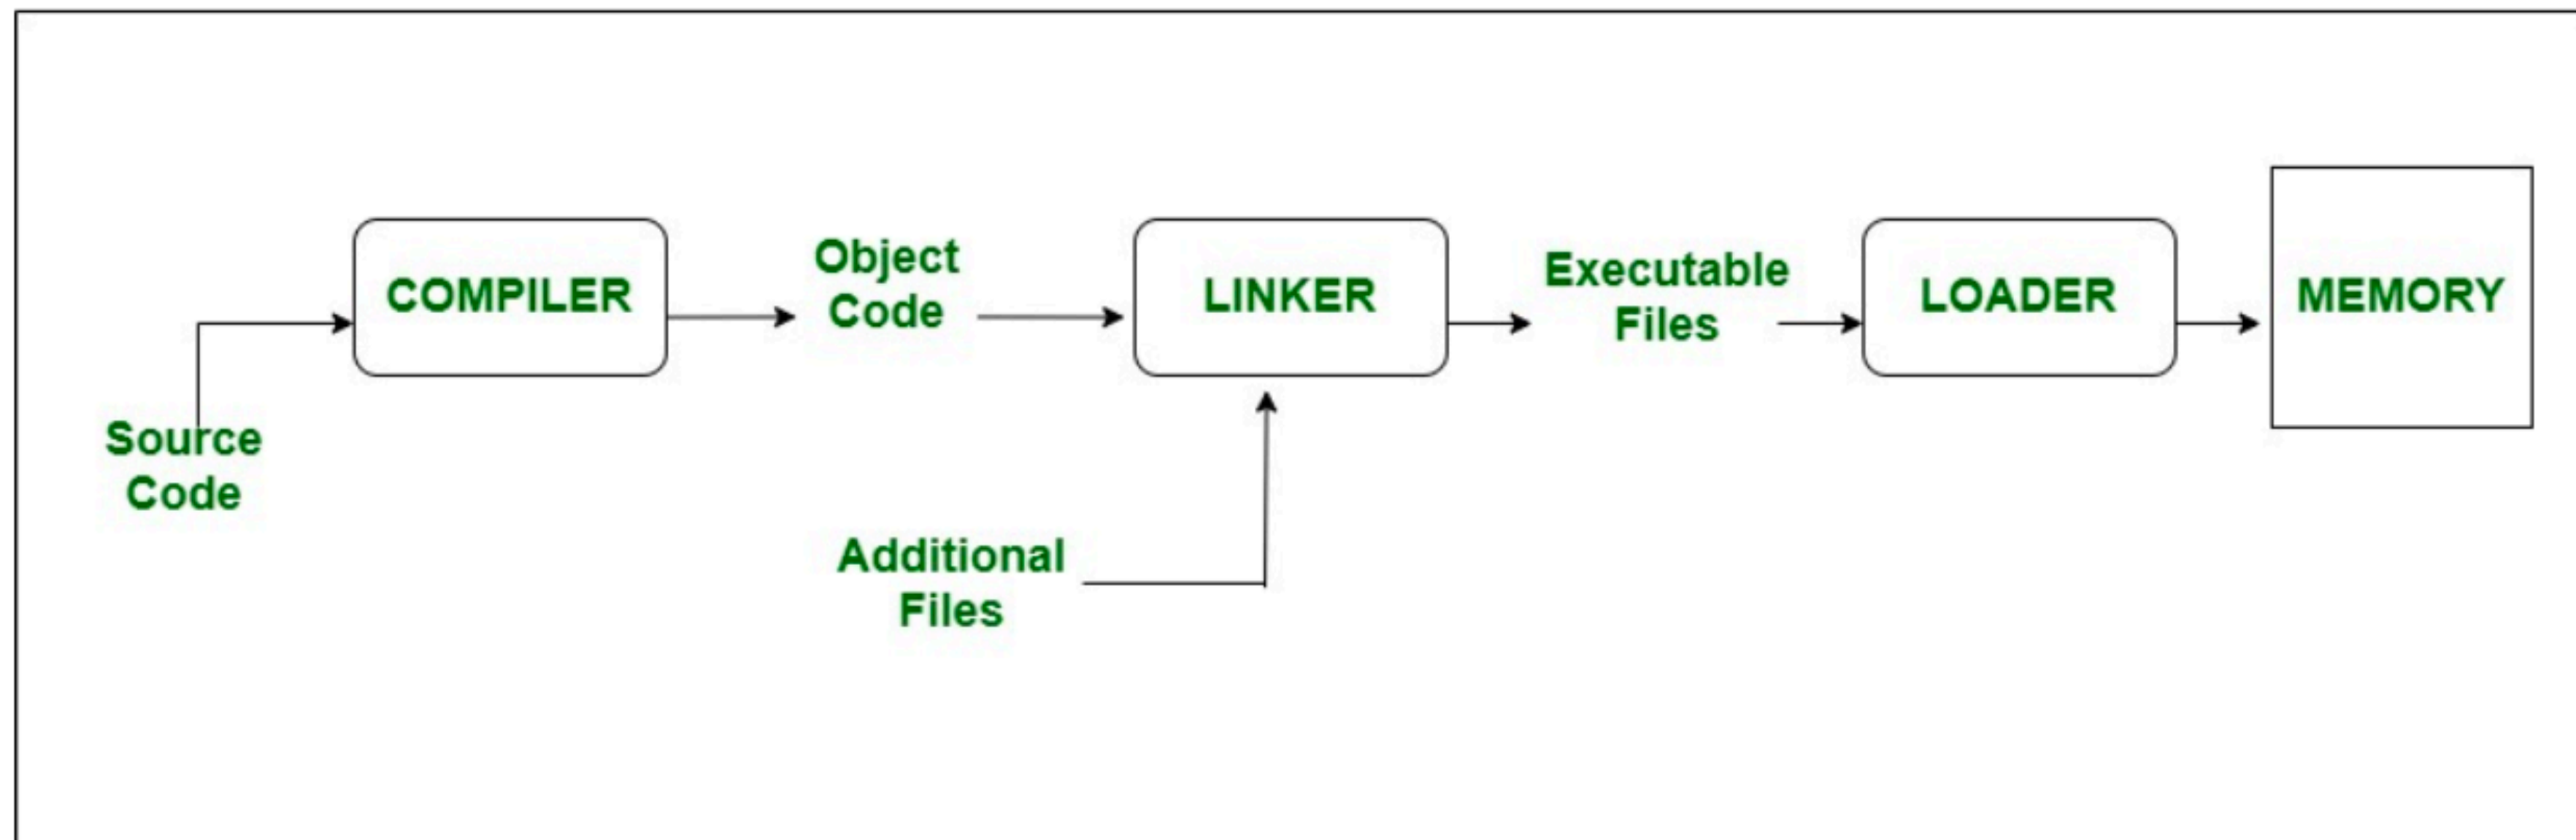
Why do we need to relocate?

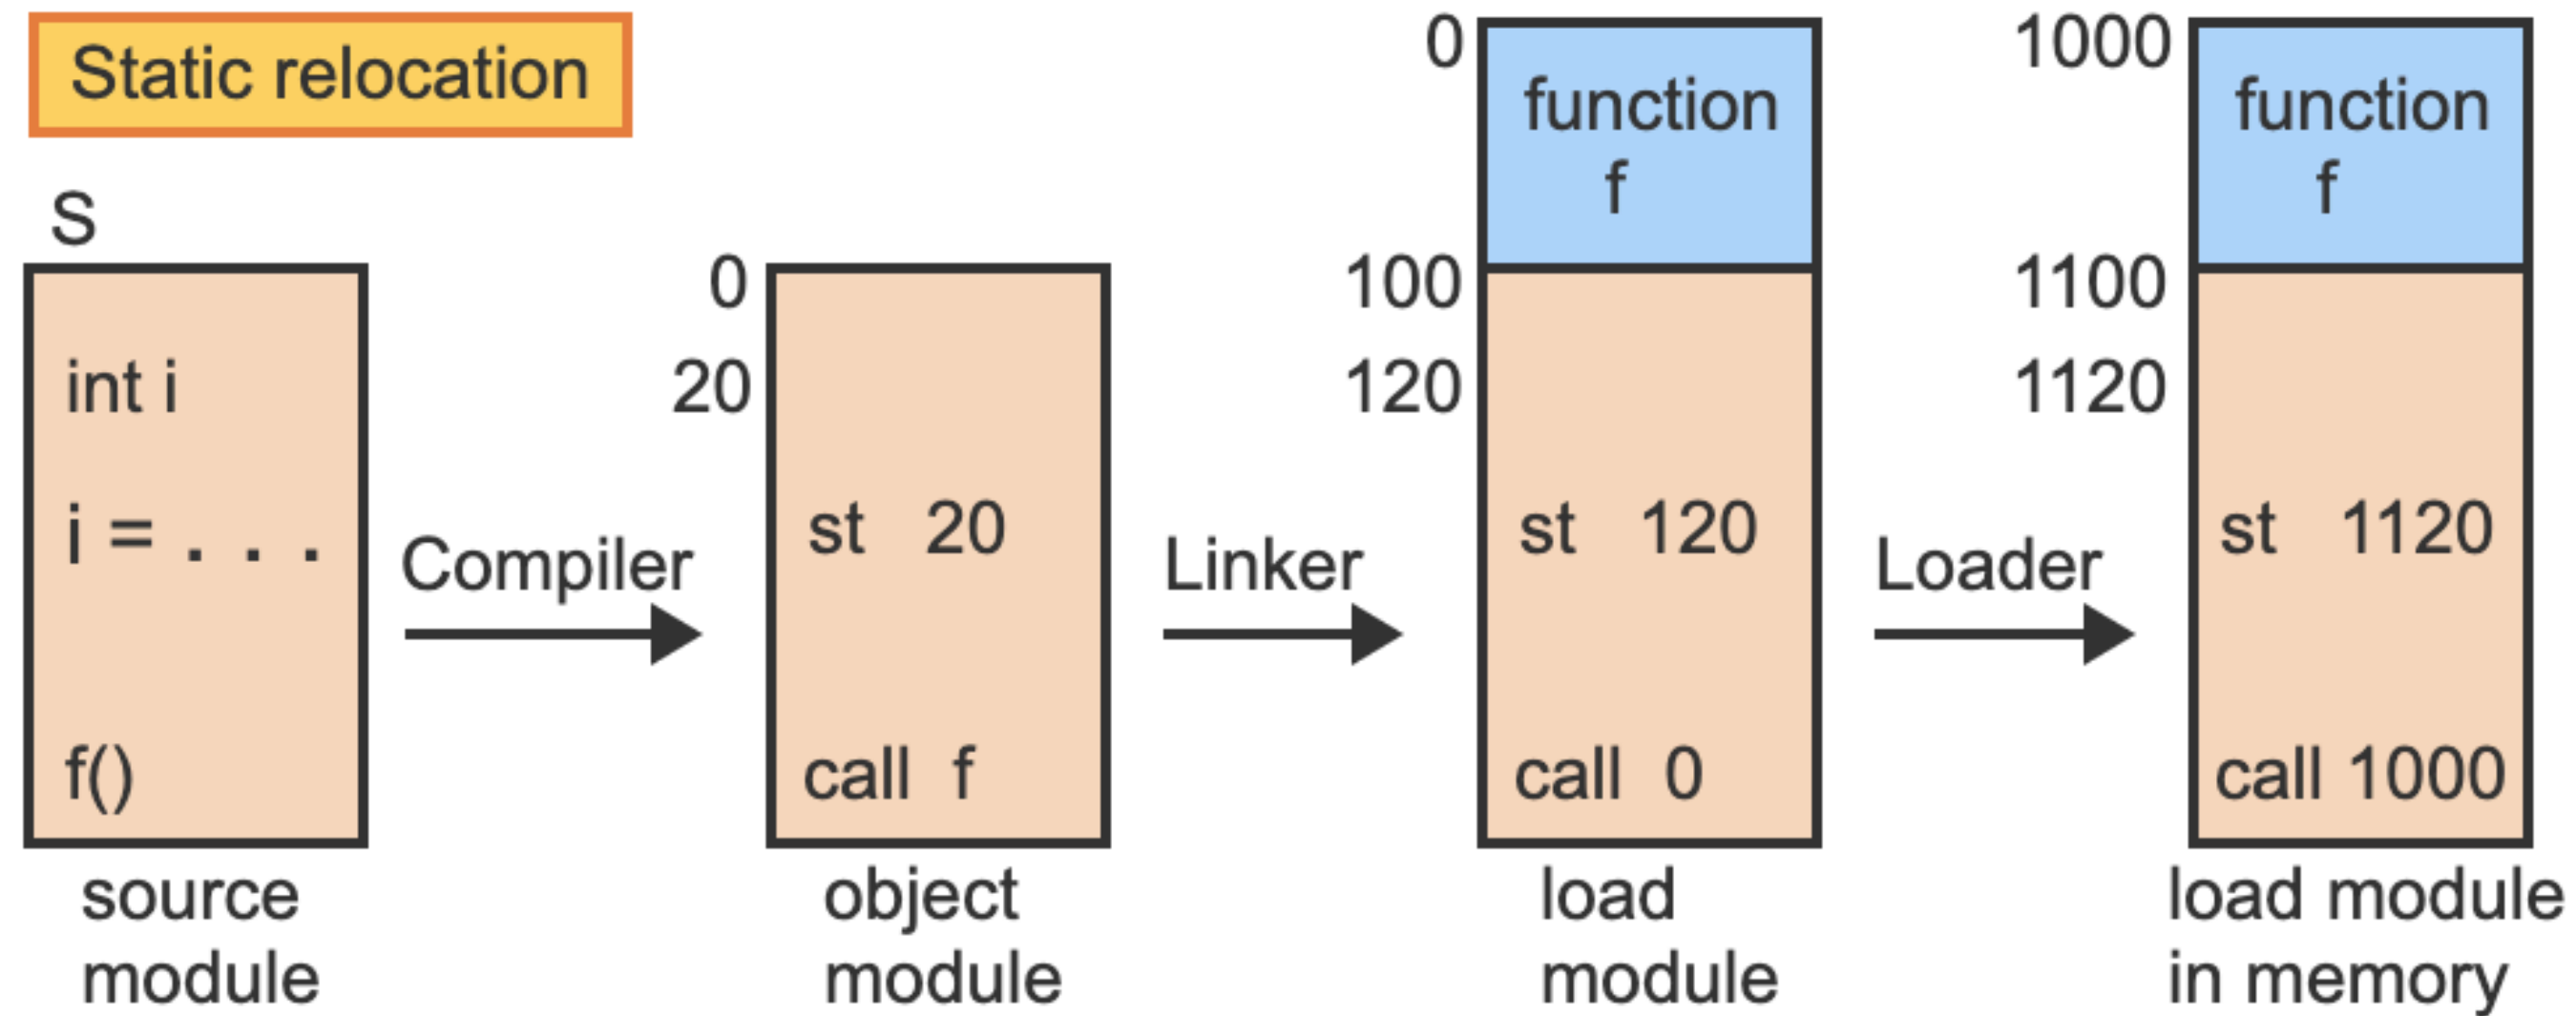
- Since the physical addresses of the program components are not known until load time,
- the compiler/assembler and the linker assume logical address spaces starting with address 0.
- During each step of the program transformation, the logical addresses of instructions and data may be changing.
- Only during the final step of program loading are all logical addresses bound to actual physical addresses in memory.

What is relocation?

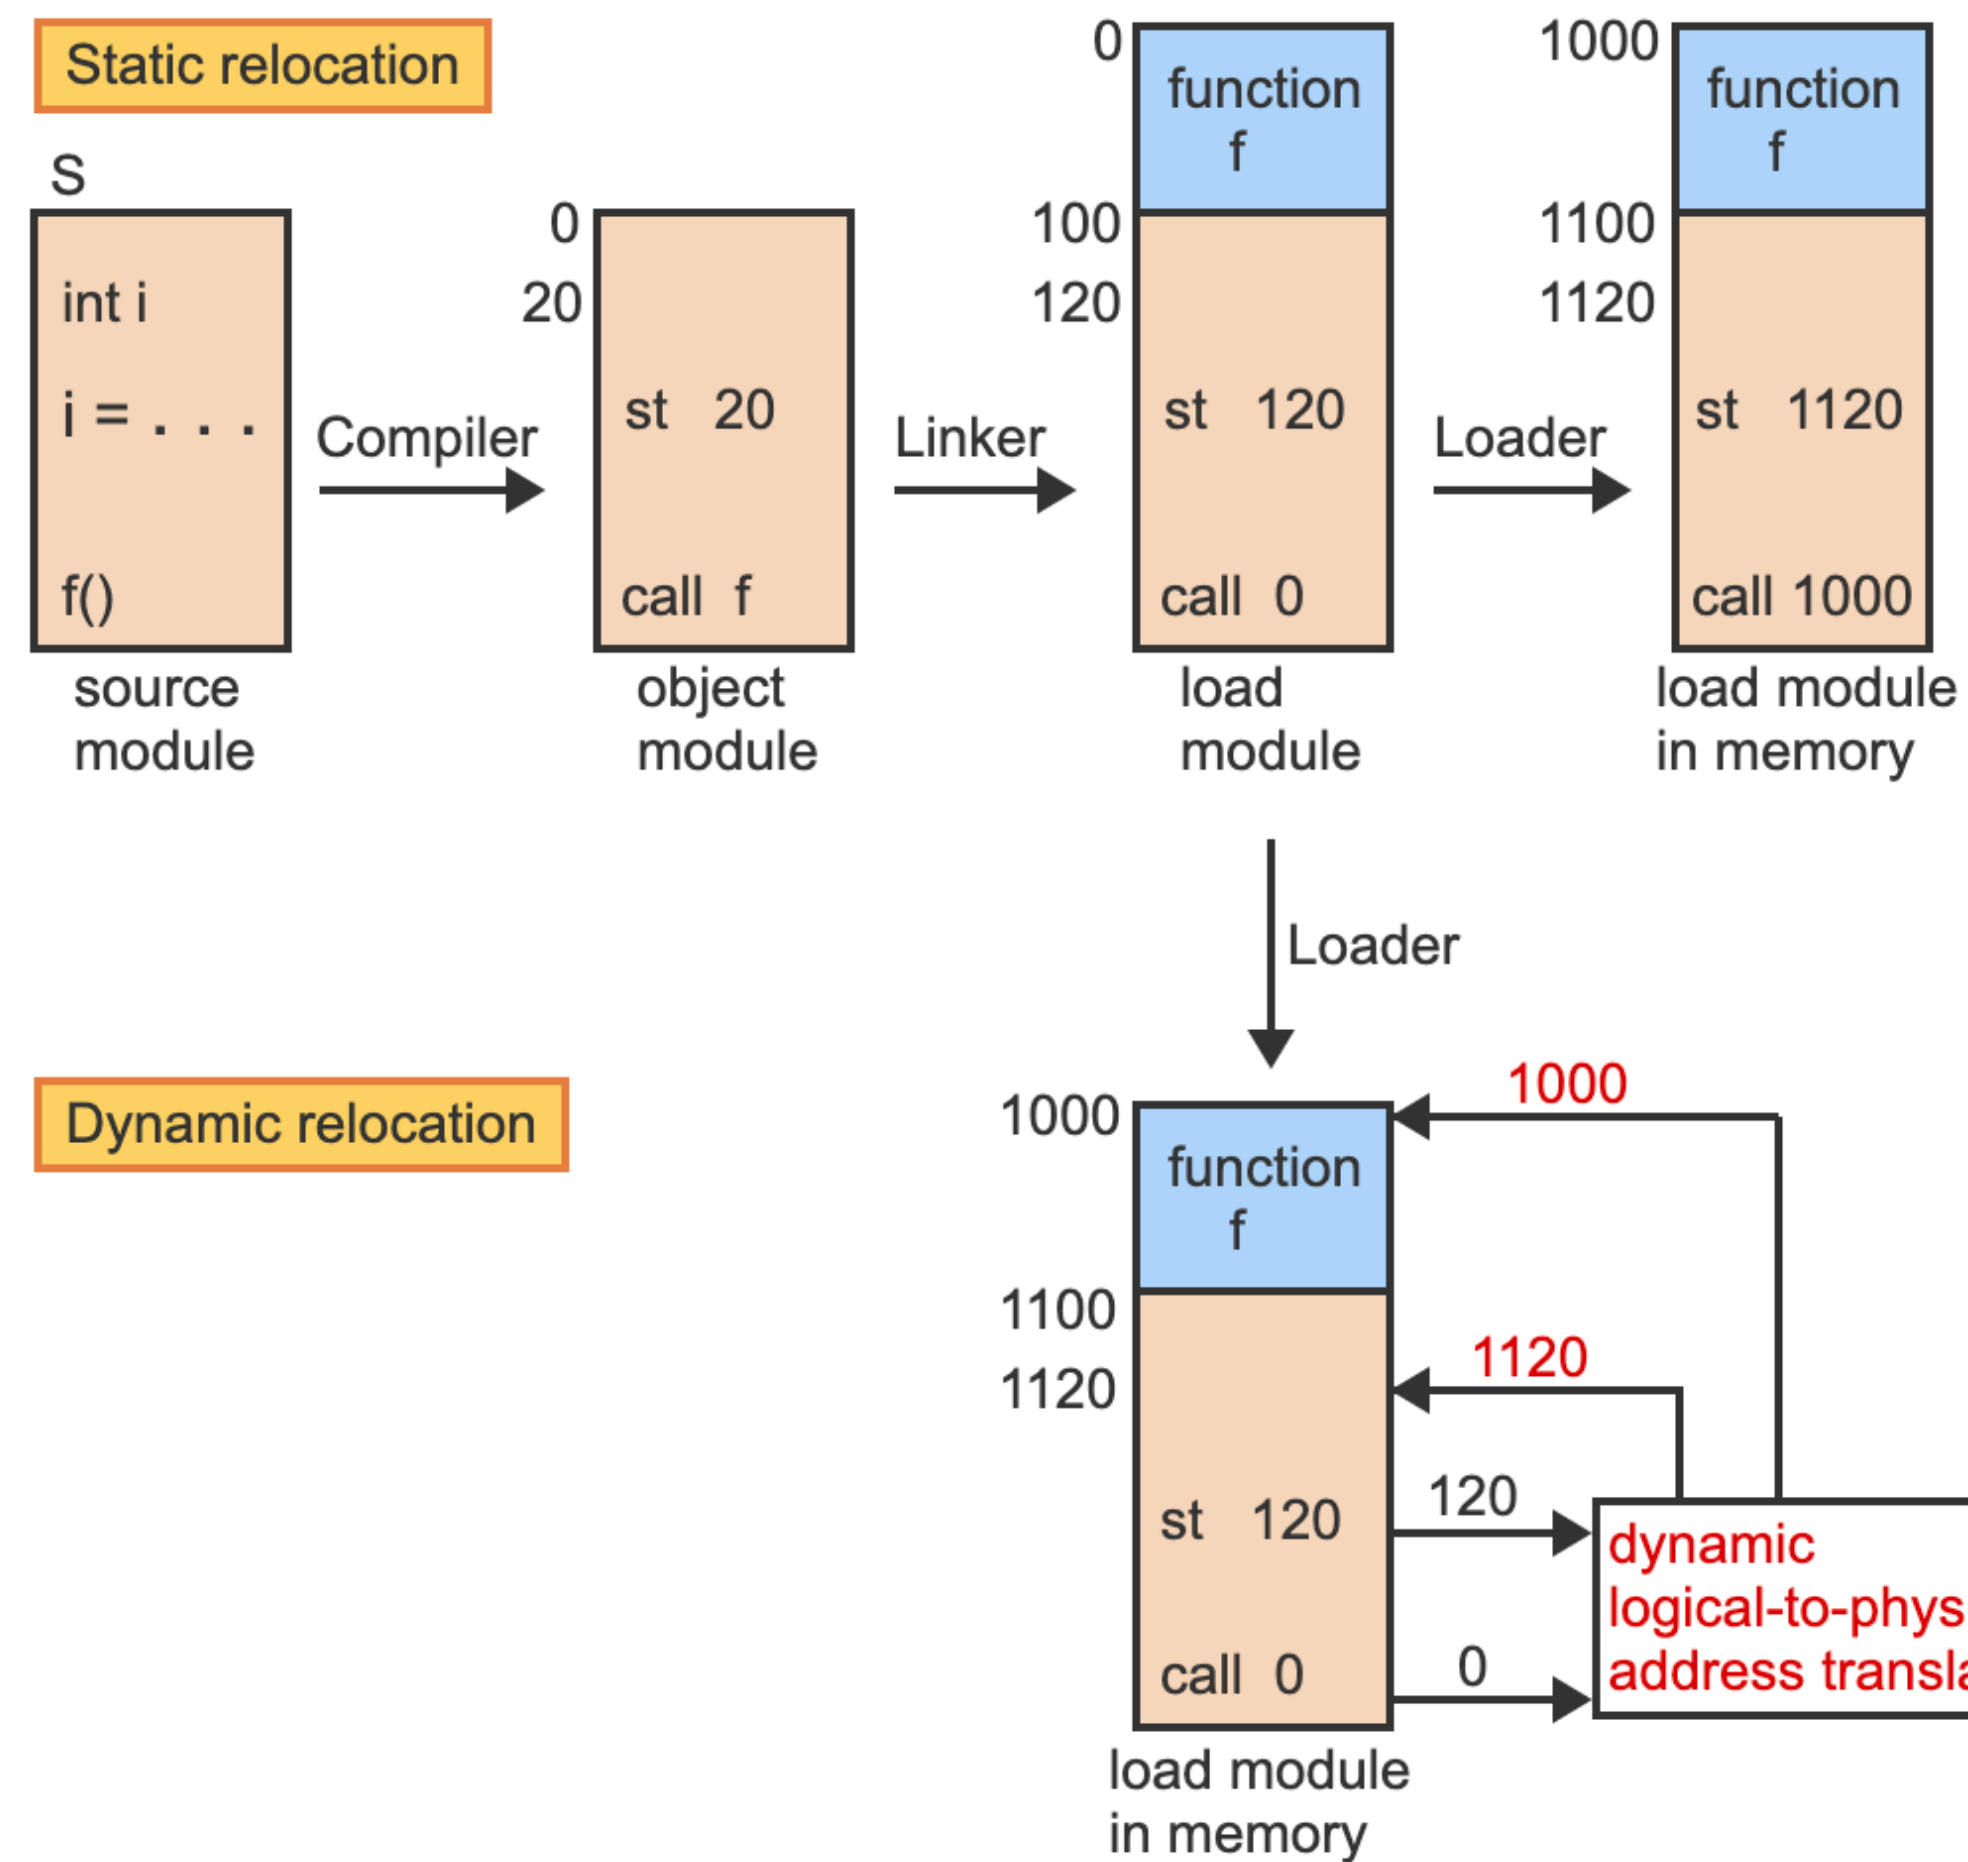
- Program relocation is the act of moving a program component from one address space to another.
- The relocation may be between two logical address spaces or from a logical address space to a physical address space.

Taking code to execution:



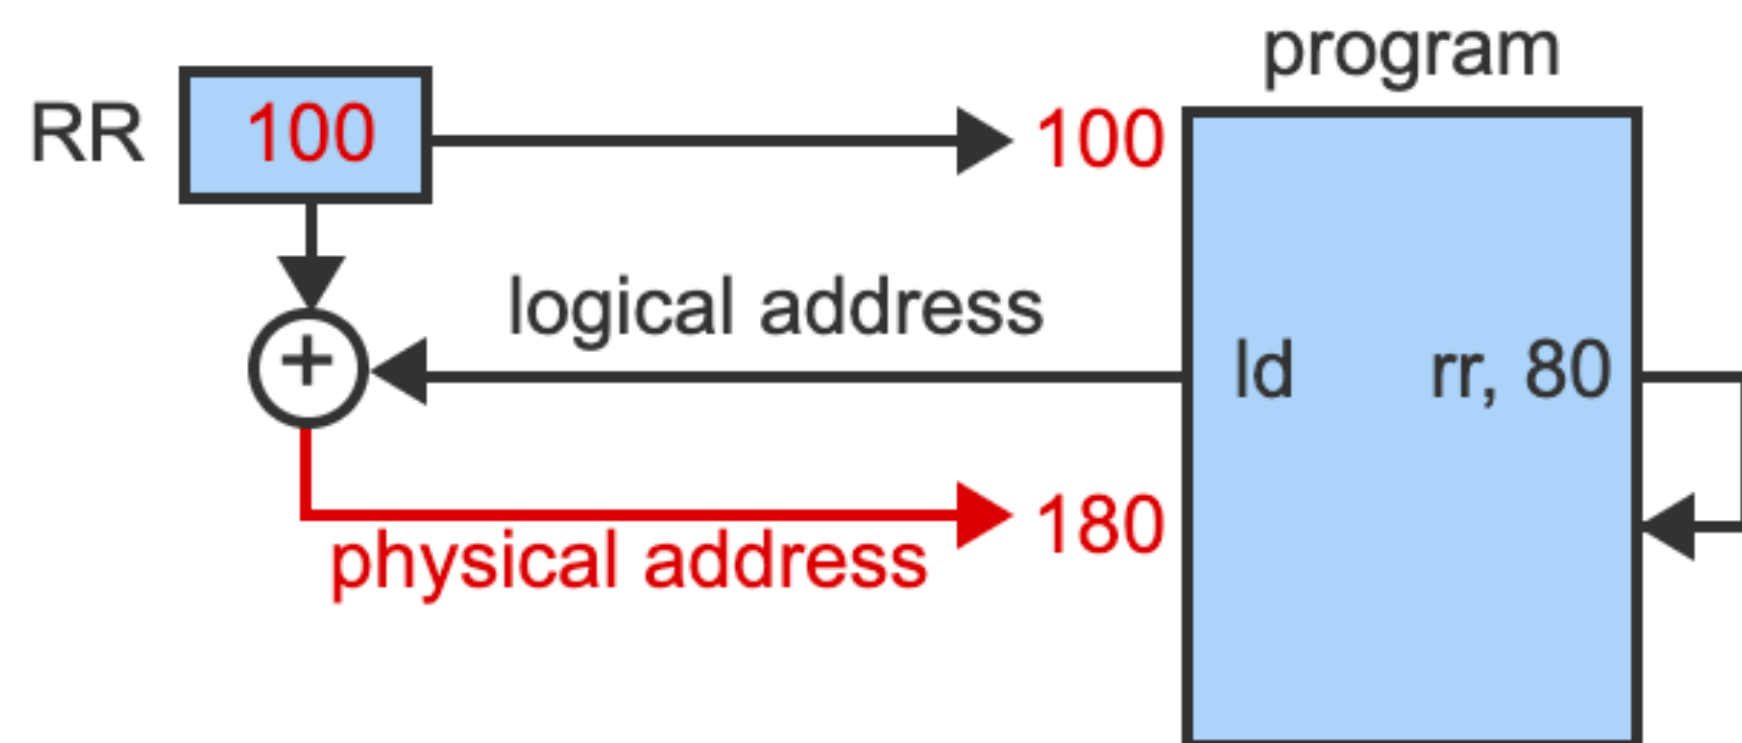
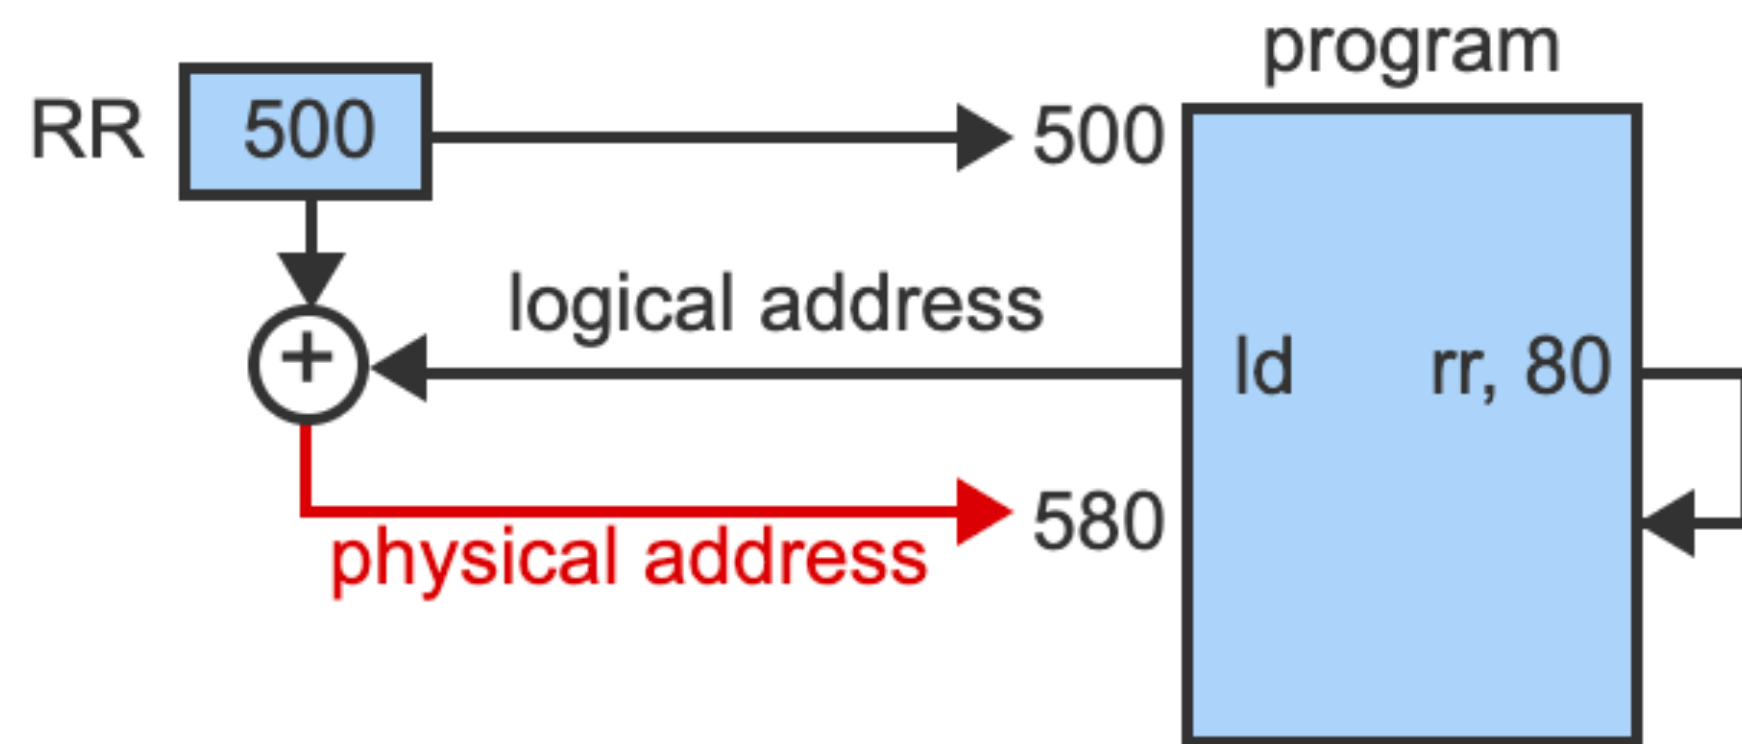


- A program consists of multiple source modules, each written in a high-level language. Module S defines a variable `i`, which is later assigned a value. Function `f` resides in a different module.
- The compiler assumes a logical address 0 and assigns `i` to address 20. The store instruction `st` then references location 20. The location of `f` is unknown and is kept as a symbolic external reference.
- The linker also assumes a logical starting address of 0 and combines all object modules into one load module. S is relocated to reflect the starting address 100. The call now refers to the start of `f`.
- **With static relocation**, the loader must adjust all logical addresses to physical addresses when copying the module into memory at the starting physical address of 1000



- With **dynamic** relocation, the load module is copied into memory also starting at address 1000 but without any changes to the addresses in the code
- Relocation is done at runtime by adding the starting address 1000 to every address when referenced

Dynamic Relocation using Relocation Register



- When a single-module program is loaded into memory, the starting address 500 is loaded into the relocation register RR as the base address for future address calculations
- The load instruction *ld* uses a flag, RR, to indicate that the operand 80 is not an absolute value but a logical address referring to 580
 - The CPU adds 500 to 80, which results in the physical address 580
- When the program is moved to a new area of memory starting at address 100, the code does not change
 - Only RR is loaded with 100, which guarantees that correct physical addresses are generated at runtime


Apple II+



- June 1979, 44 years ago for \$1200 (\$5000 today)
- 16KB, 32KB, 48KB of RAM expandable to 64KB using an expansion card

Single User Contiguous Scheme

- Program was loaded into memory contiguously
- If Program was too big it would not run unless memory was increased.
 - Memory was “expensive back then (\$6000 in today’s money) for 16KB.



Does it fit?

Operating System	10K
Program 1 (40K)	50K
Unused Main Memory	

- If it doesn't fit, the Program could not load or run. The programs of "yesterday" had no concept of Virtual Memory (as everything contained in the program)

Advantages

- Single User Continuous Scheme Advantages:
 - Jobs are allocated sequentially
 - i.e. Memory Manager is simple
 - i.e. OS has to do less
 - OS checks if there's room
 - If not, job aborted

Disadvantages

- Single User Continuous Scheme Disadvantages:
 - Only one program at a time
 - Memory cannot be shared
 - Memory wasted
 - If program is small, extra memory lies there doing nothing

How to fix it?

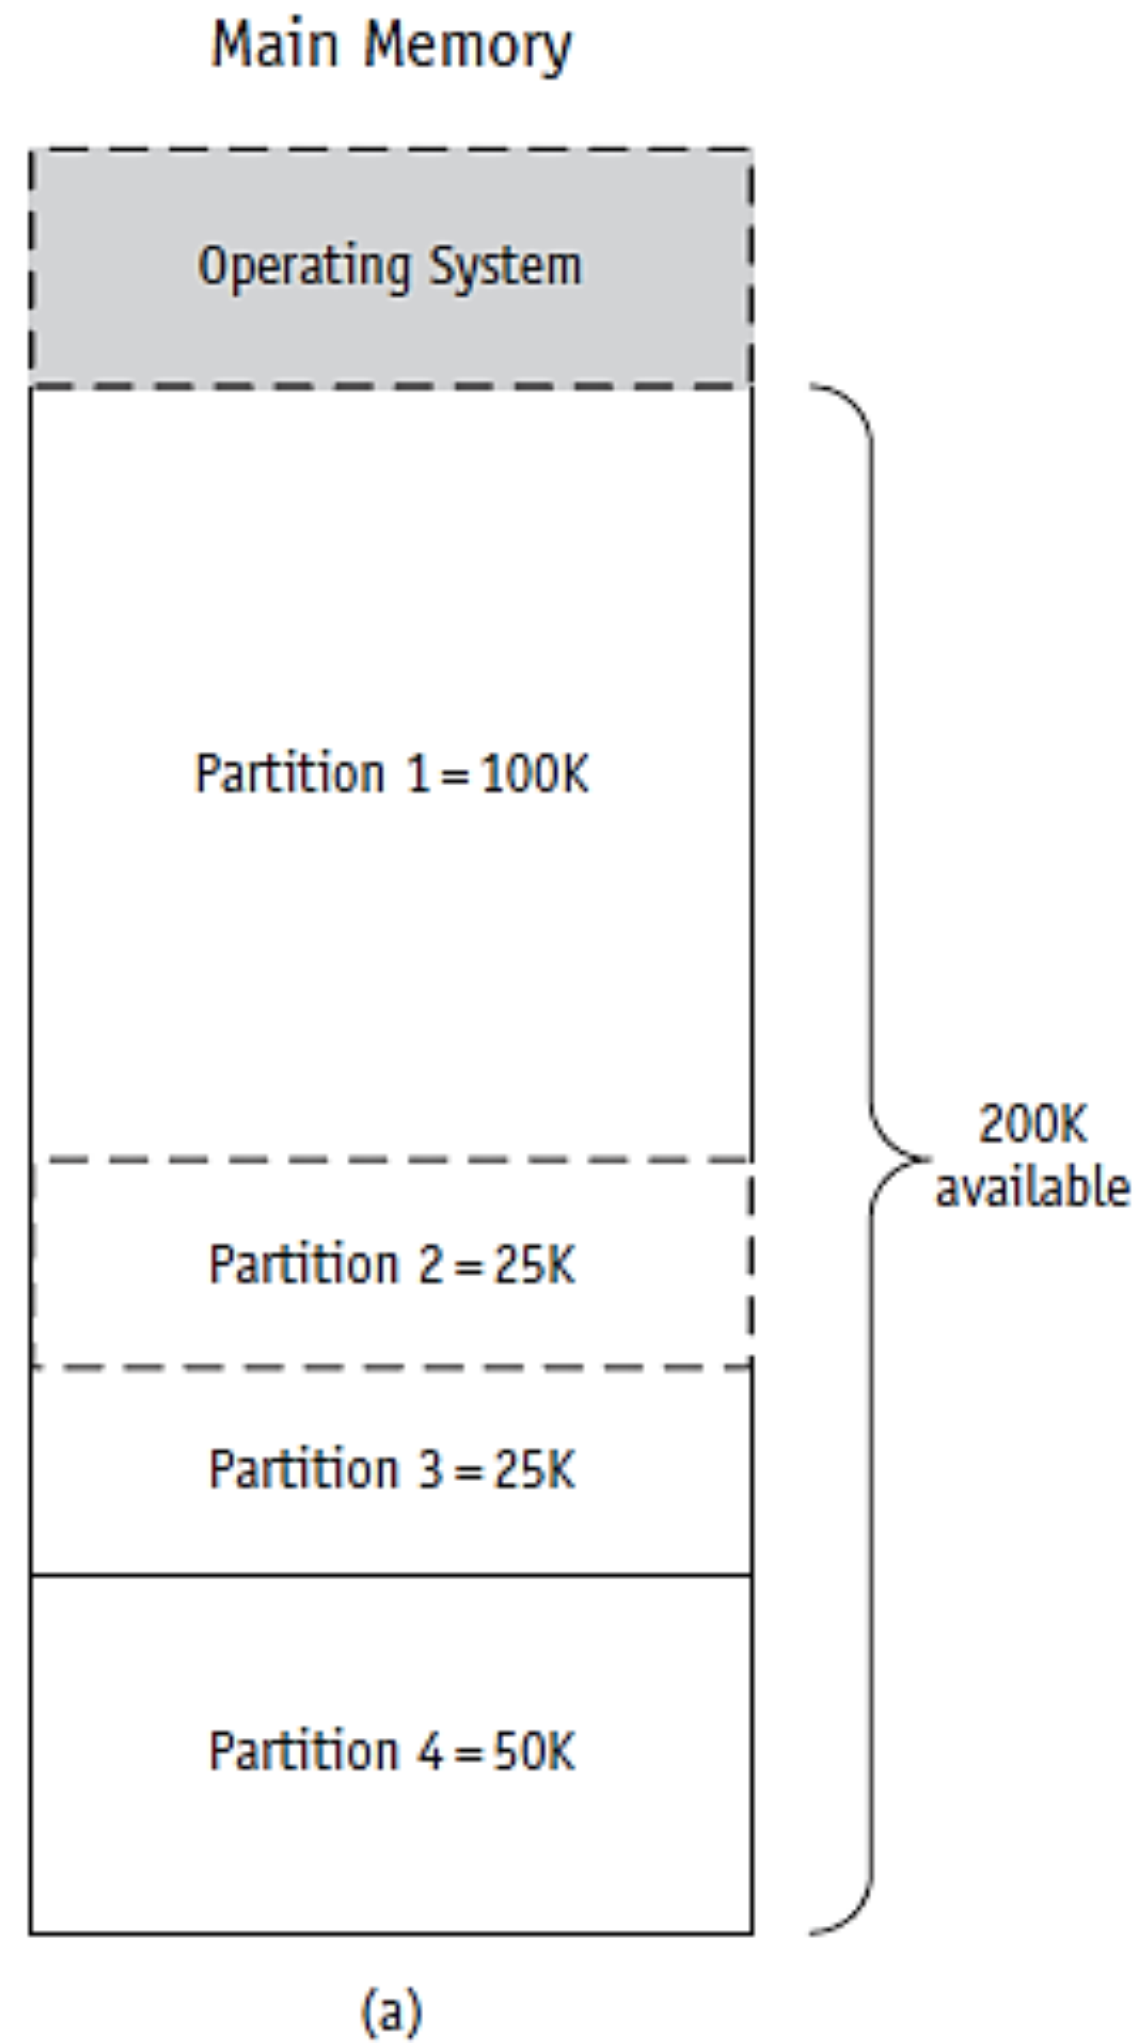
- Fixed Partitions!
- Sys Admins would split memory into partitions
 - Not necessarily same size
 - Probably in accordance to programs that ran
- e.g. 64K -> 1x32K partitions & 1x16K partitions & 2x8K partitions

How to fix it?

- Fixed Partitions supported multiple jobs

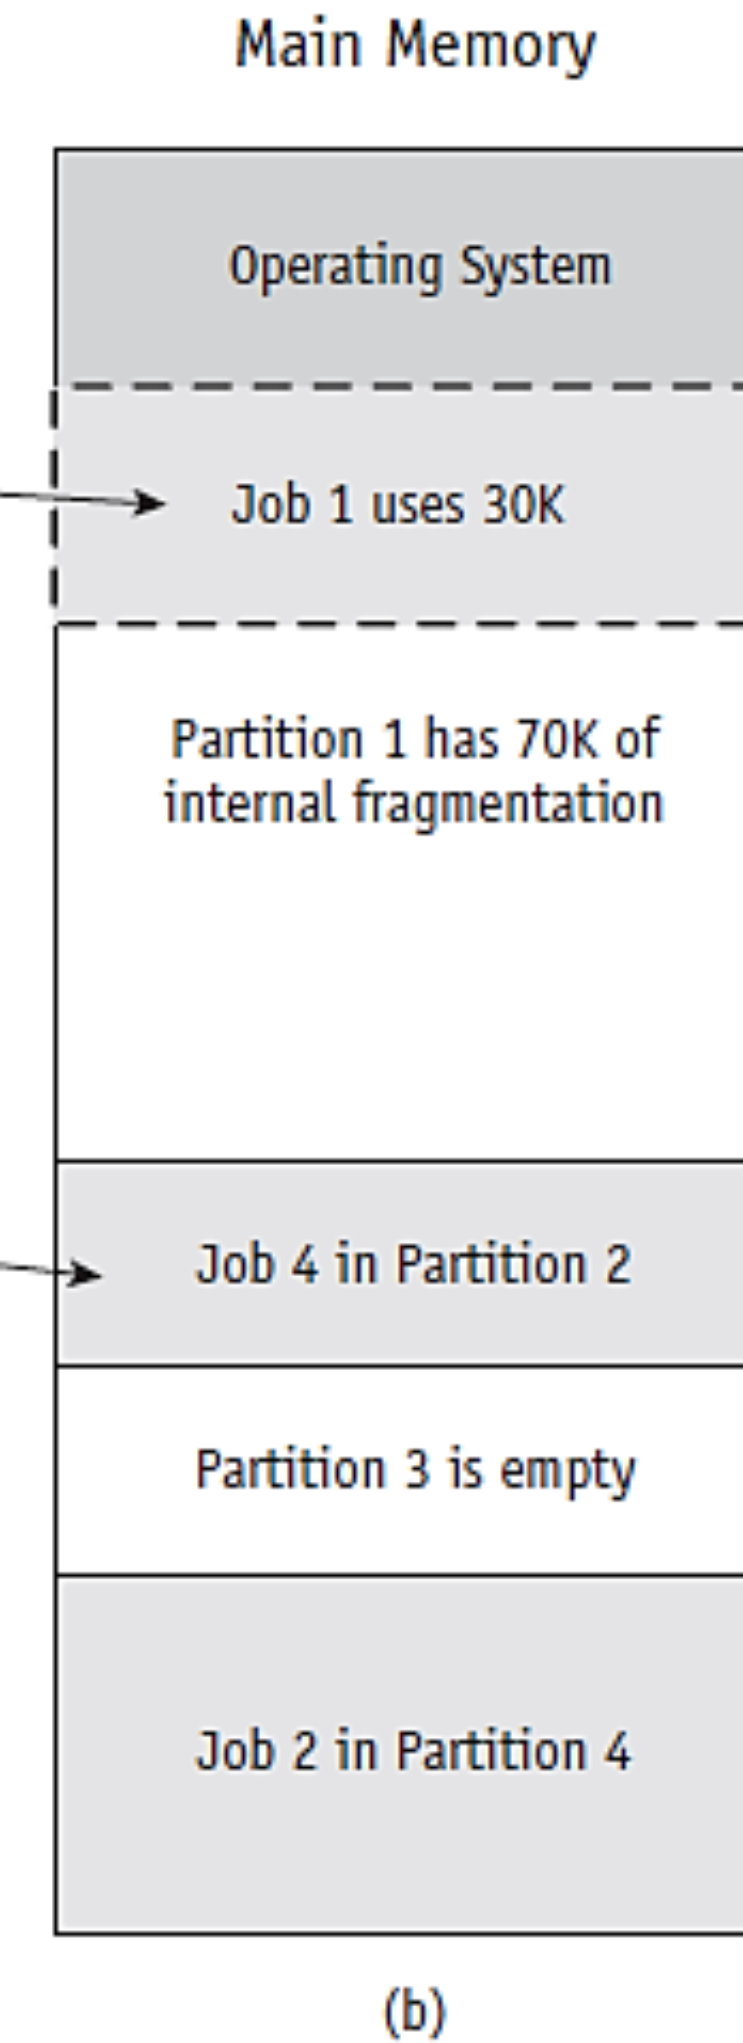
Partition Size	Memory Address	Access	Partition Status
100K	200K	Job 1	Busy
25K	300K	Job 4	Busy
25K	325K		Free
50K	350K	Job 2	Busy

Job 3 is starved



JOB LIST :

Job 1 = 30K
Job 2 = 50K
Job 3 = 30K (waiting)
Job 4 = 25K



Advantages/ Disadvantages

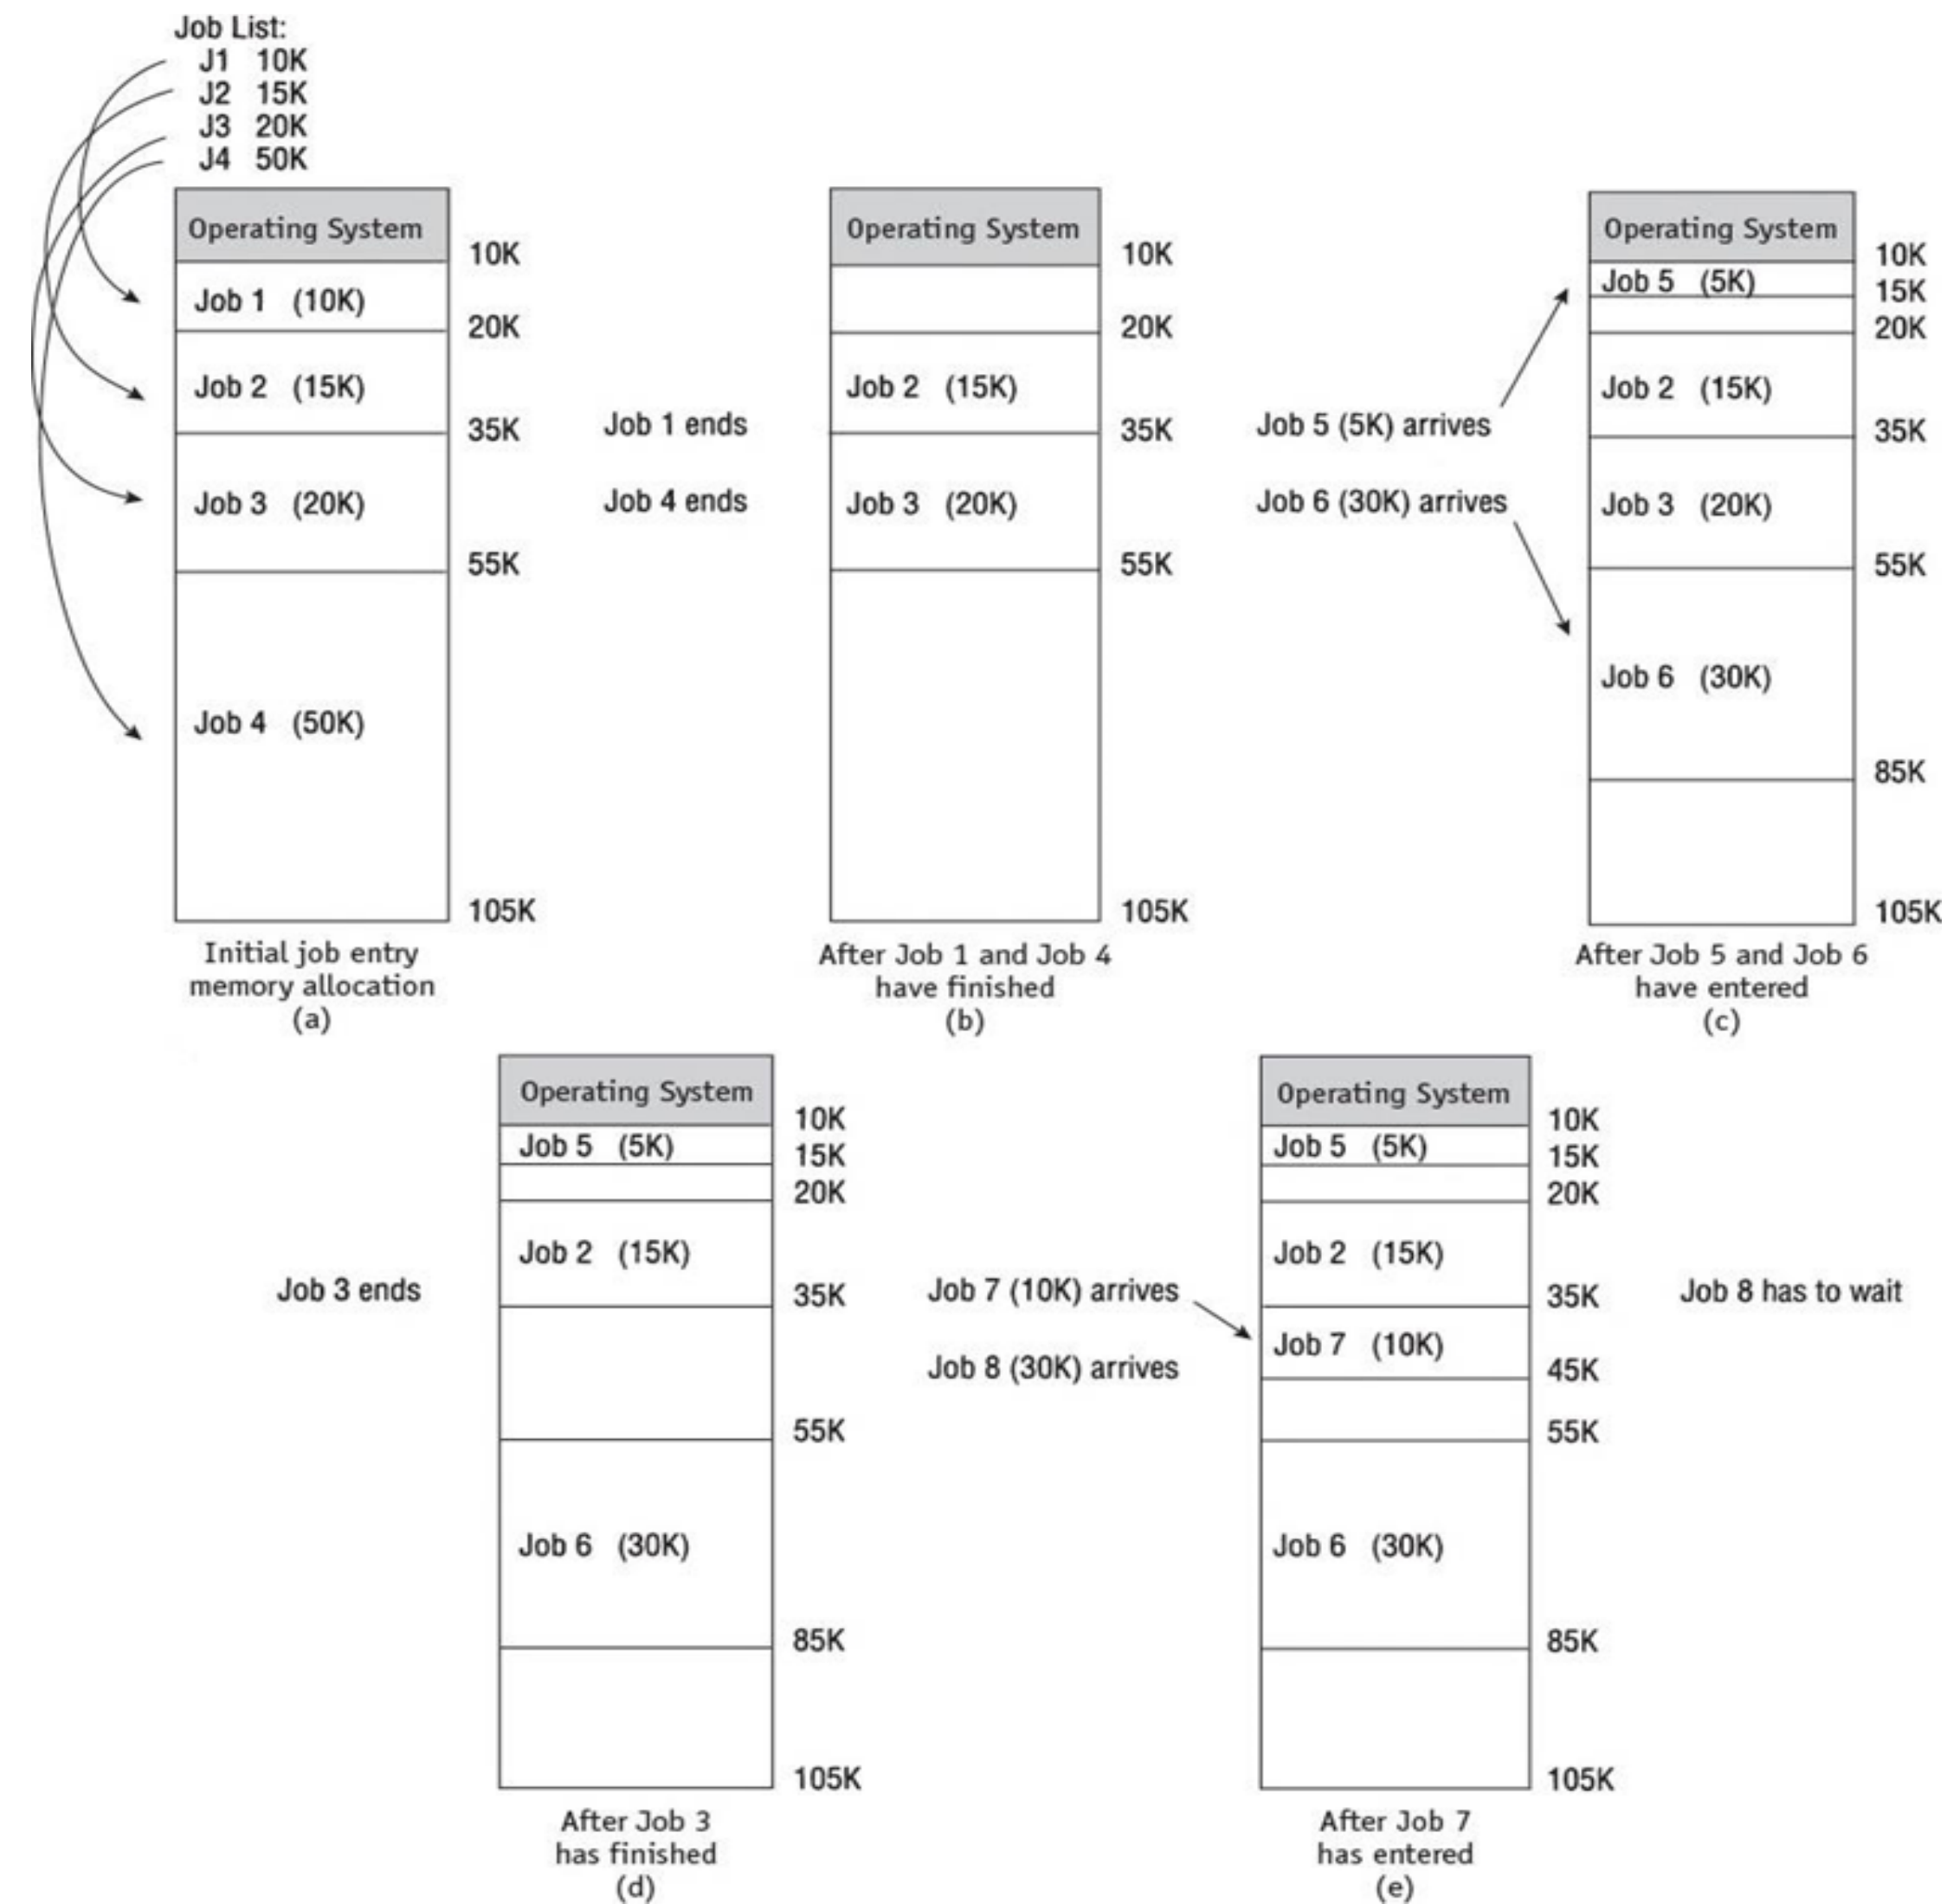
- Fixed Partitions:
 - Advantages
 - Support multi-programming
 - Protected with interfering from each other
 - Disadvantages
 - Partition size may not reflect reality
 - Sysadmin has to constantly keep updating sizes

Is dynamic better?

Dynamic Partitions


- Main Memory: Partitioned
 - One contiguous partition per job
- Advantages:
 - Better use of memory
- Disadvantages:
 - De-Fragmentation could result into more management

Dynamic Partitions



Internal vs External Fragmentation

My Living Room



First Fit Allocation

- Using a first-fit scheme,
 - Job 1 claims the first available space
 - Job 2 then claims the first partition large enough to accommodate it,
 - but by doing so it takes the last block large enough to accommodate Job 3
 - Therefore, Job 3 must wait until a large block becomes available,
 - even though there's 75K of unused memory space (internal fragmentation)
- Let's see..

First Fit Allocation

Job List:

Job number	Memory requested
J1	10K
J2	20K
J3	30K*
J4	10K

Memory List:

Memory location	Memory block size	Job number	Job size	Status	Internal fragmentation
10240	30K	J1	10K	Busy	20K
40960	15K	J4	10K	Busy	5K
56320	50K	J2	20K	Busy	30K
107520	20K			Free	
Total Available:	115K	Total Used:	40K		

Best Fit Allocation

Job List:

Job number	Memory requested
J1	10K
J2	20K
J3	30K
J4	10K

Memory List:

Memory location	Memory block size	Job number	Job size	Status	Internal fragmentation
40960	15K	J1	10K	Busy	5K
107520	20K	J2	20K	Busy	None
10240	30K	J3	30K	Busy	None
56320	50K	J4	10K	Busy	40K
Total Available:	115K	Total Used:	70K		

Next Fit

- Starts searching from last allocated block for next available block

Worst Fit

- Worst-fit: allocates largest free available block
- Opposite of best-fit
- Why would it be useful though? Why is it even a strategy?
-

Worst Fit

- Worst-fit: allocates largest free available block
- Opposite of best-fit
- Why would it be useful though? Why is it even a strategy?
 - Produces largest possible fragmentation!

Best Strategy?

- Simulations have shown that both first fit and best fit are better than worst fit in terms of decreasing time and storage utilization.
- Neither first fit nor best fit is clearly better than the other in terms of storage utilization, but first fit is generally faster.

Deallocation

- For Fixed Partition System, its simple as status goes from busy to free
- Dynamic is much harder as it tries to combine free areas together and make one free area

Deallocation

- 3 cases:
 - When block to be deallocated is adjacent to a free block
 - When block is between two free blocks
 - When block to be deallocated is isolated from other blocks

Joining two free blocks

Beginning Address	Memory Block Size	Status
4075	105	Free
5225	5	Free
6785	600	Free
7560	20	Free
(7600)	(200)	(Busy) ¹
*7800	5	Free
10250	4050	Free
15125	230	Free
24500	1000	Free

What do you think will happen after 7600 changes to free?

If you think that it will be 7600 with a block size of 205, you are correct

Joining three free blocks

Beginning Address	Memory Block Size	Status
4075	105	Free
5225	5	Free
6785	600	Free
*7560	20	Free
(7580)	(20)	(Busy) ¹
*7600	205	Free
10250	4050	Free
15125	230	Free
24500	1000	Free

THINK CONTIGIOUS

We will now have a Contiguous region at 7560 of 245 bytes

50% Rule

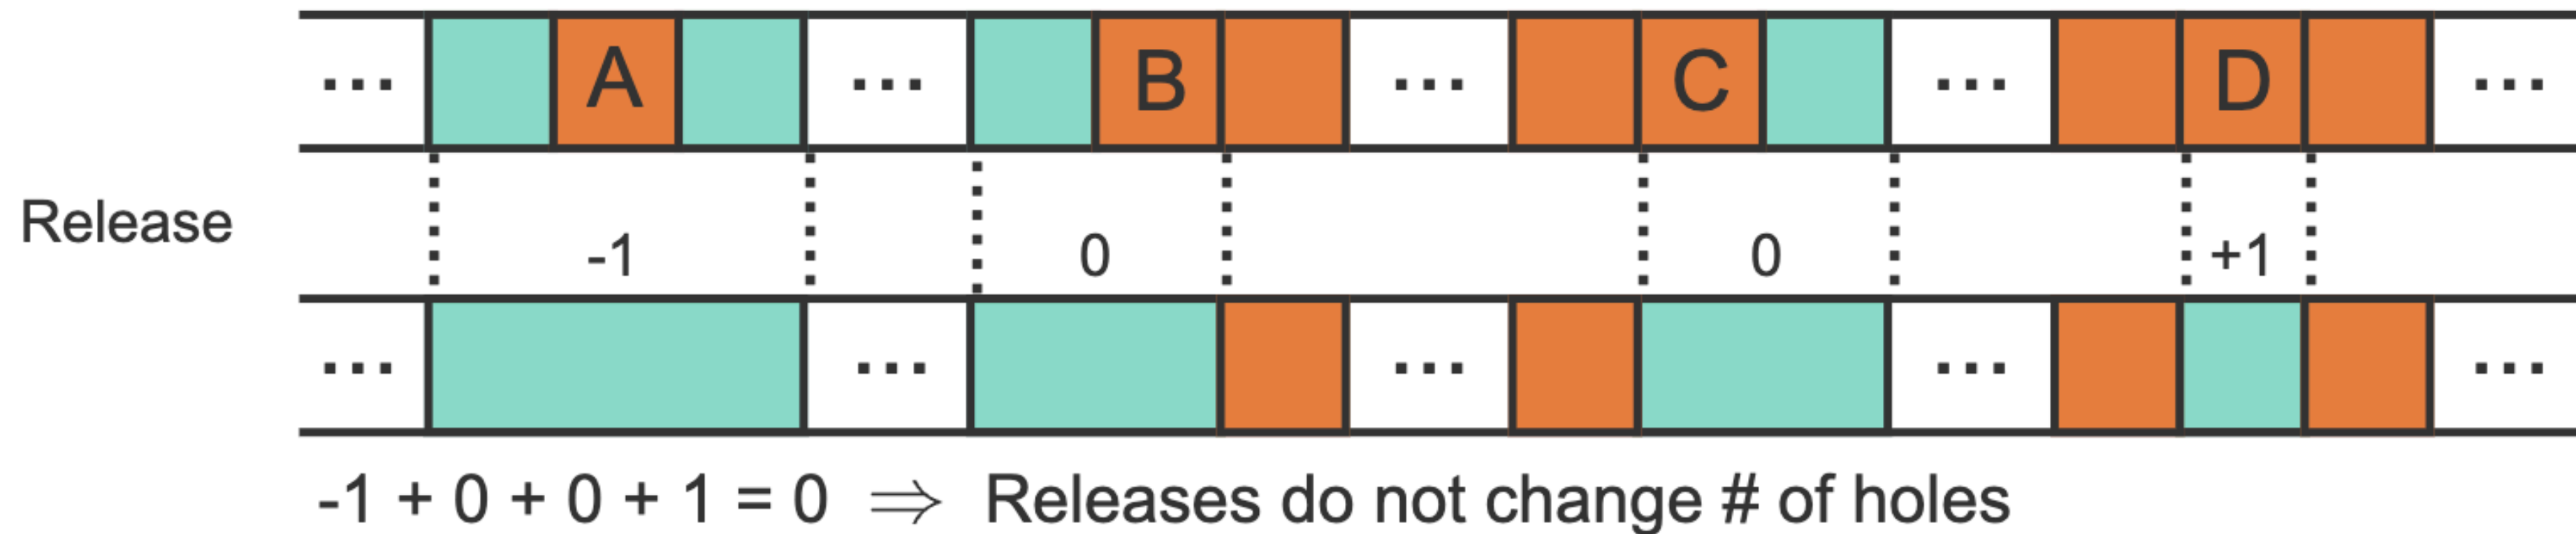
50% Rule

- If the probability of finding an exact match for a request approaches 0, one third of all memory partitions are holes and two thirds are occupied blocks
- Formally, $n = 0.5 m$, where n is the number of holes and m is the number of occupied blocks



- The 50% rule states that, on average, one hole exists for every 2 occupied blocks. The ratio cannot be changed by any allocation strategy, regardless of the block or hole sizes.

50% Rule Proof:



- 4 types of occupied blocks co-exist in memory:
- A has 2 hole neighbors. B and C have 1 hole neighbor each. D has no hole neighbors.
 - Releasing a block of type A decreases the number of holes by 1 by coalescing the released block with the 2 neighboring holes
 - Releasing a block of type B or C does not change the number of holes. Only the size of the existing hole is increased
 - Releasing a block of type D increases the number of holes by 1 by creating a new hole

Relocatable Dynamic Partitions

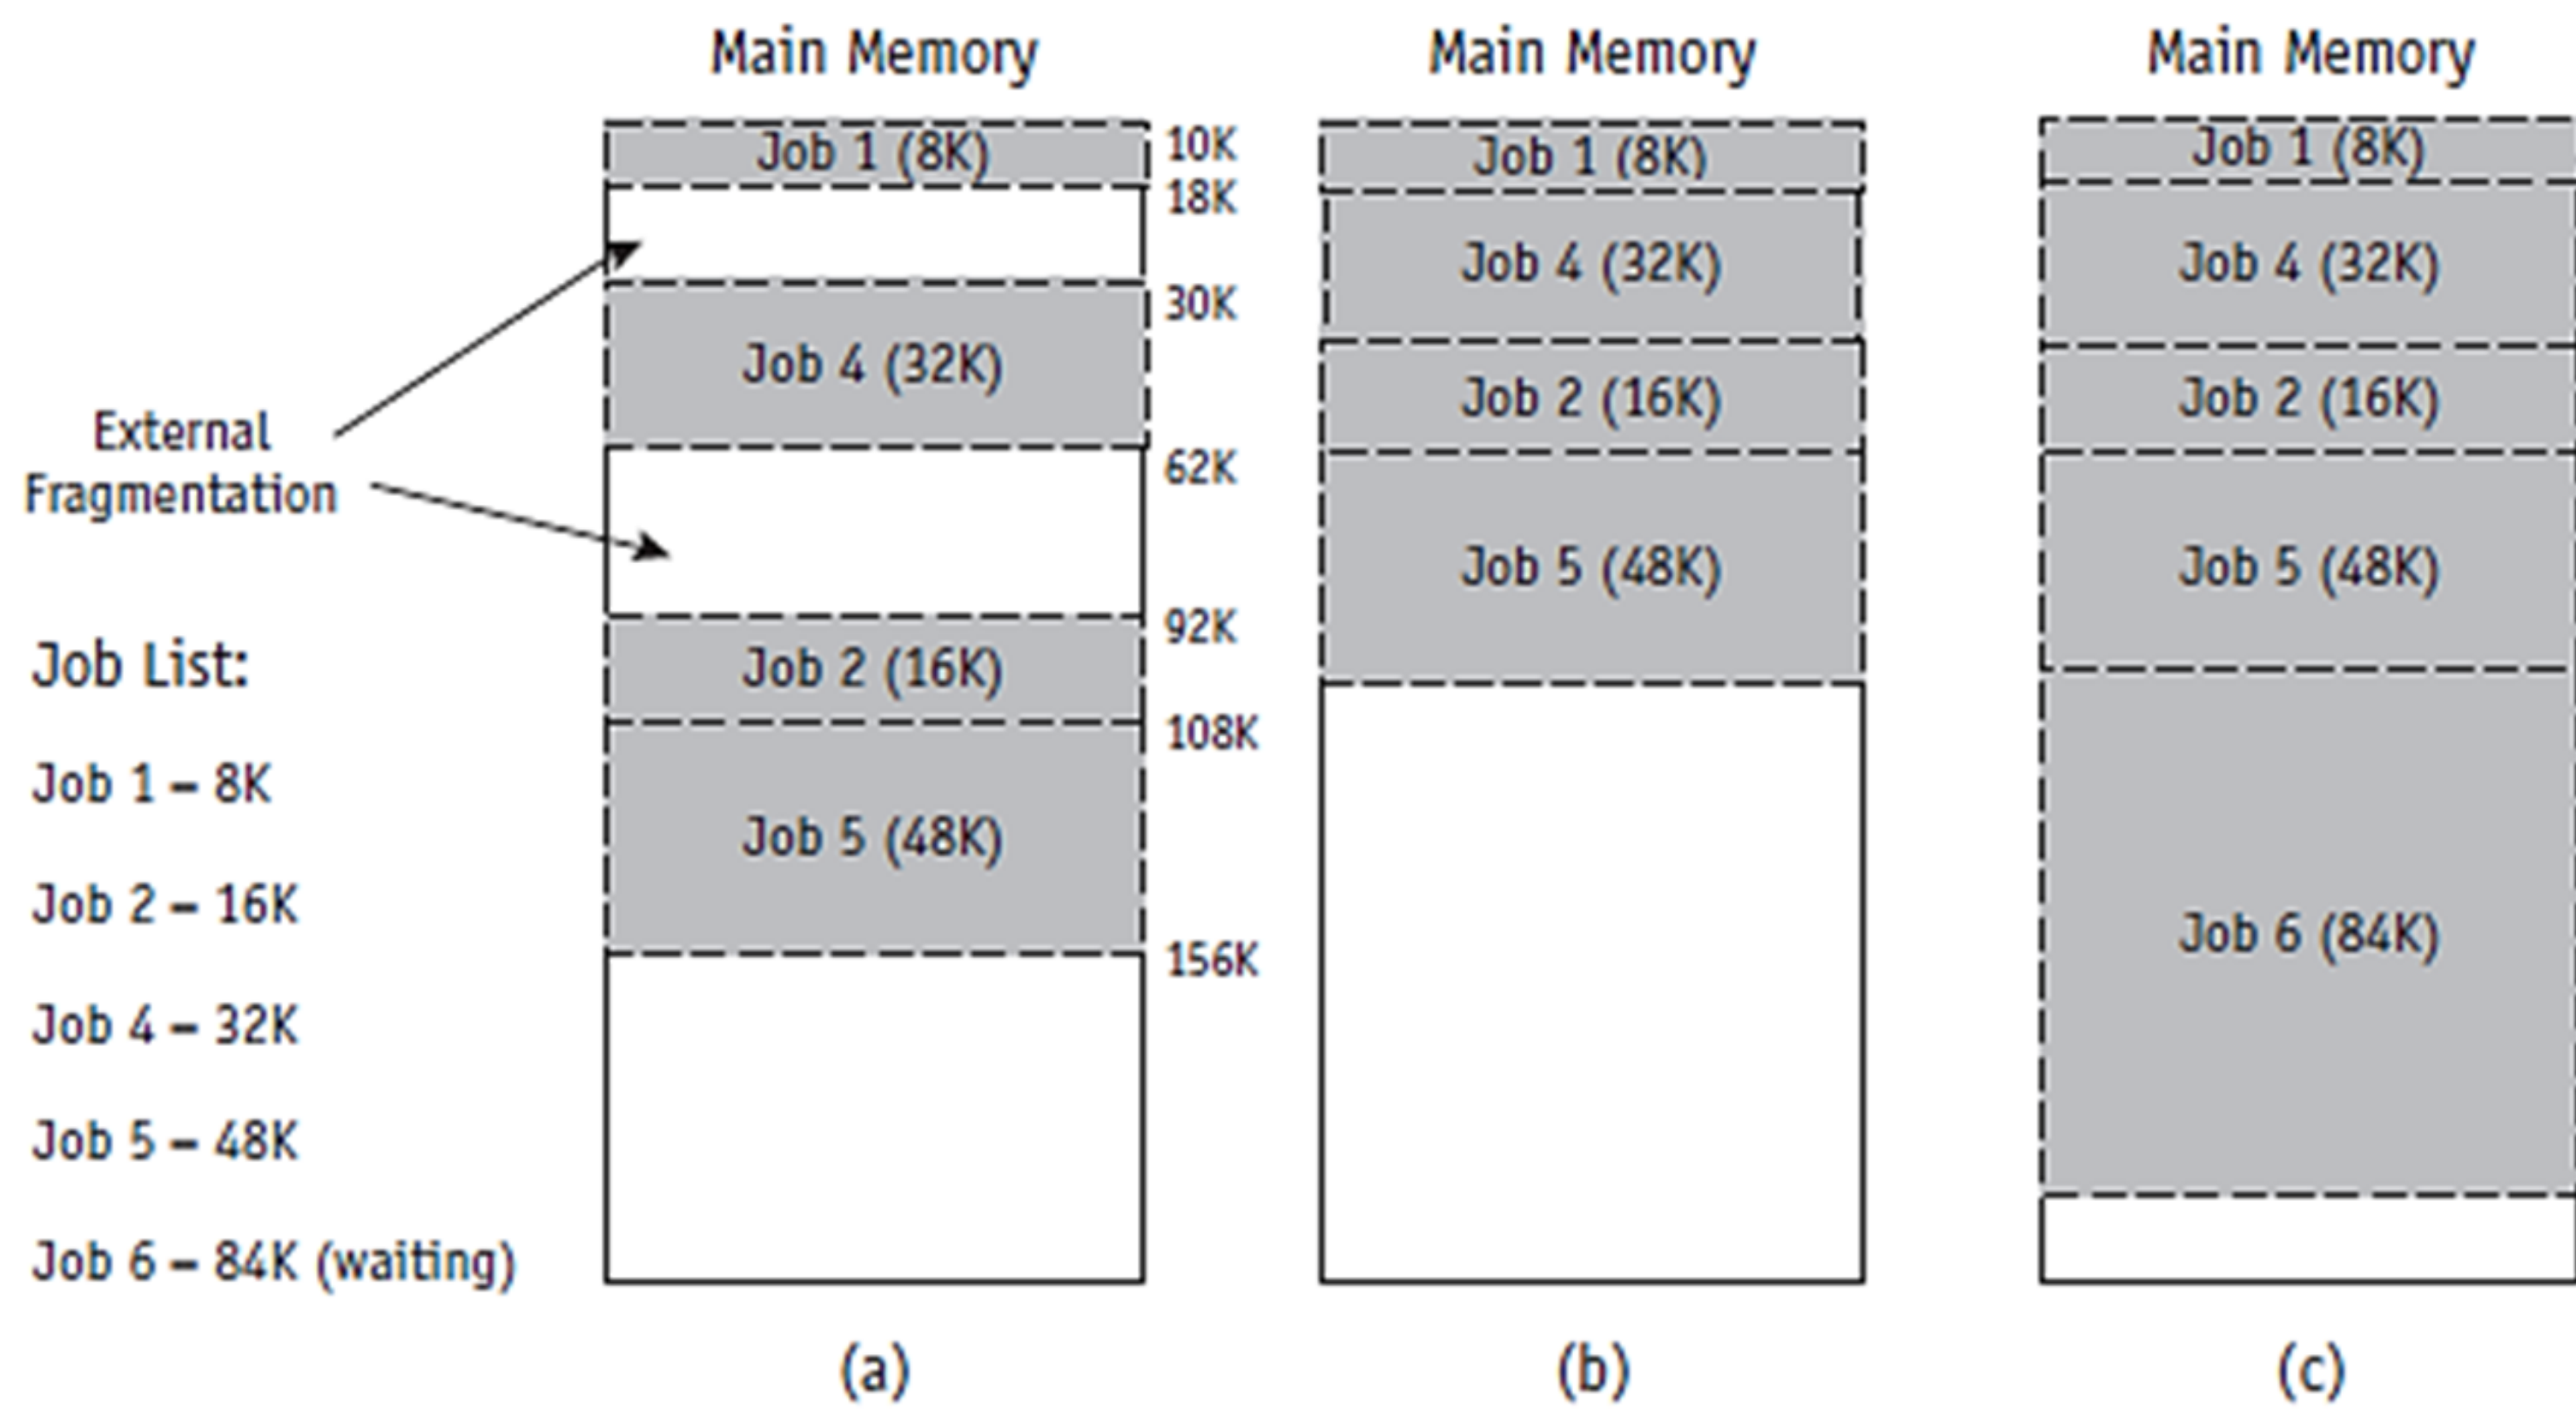
- What if the OS Memory Manager could relocate blocks in order to eliminate fragmentation.
- Concept is similar to disk defragmentation

Fragmented memory before compaction



Memory after compaction





Advantages of Compaction

- Compacting and relocating: optimizes memory use
 - Improves throughput
- Compaction: overhead
- Compaction timing options
 - When a certain memory percentage is busy
 - When there are waiting jobs
 - After a prescribed time period has elapsed
- Goal: optimize processing time and memory use; keep overhead as low as possible

Swapping vs Compaction

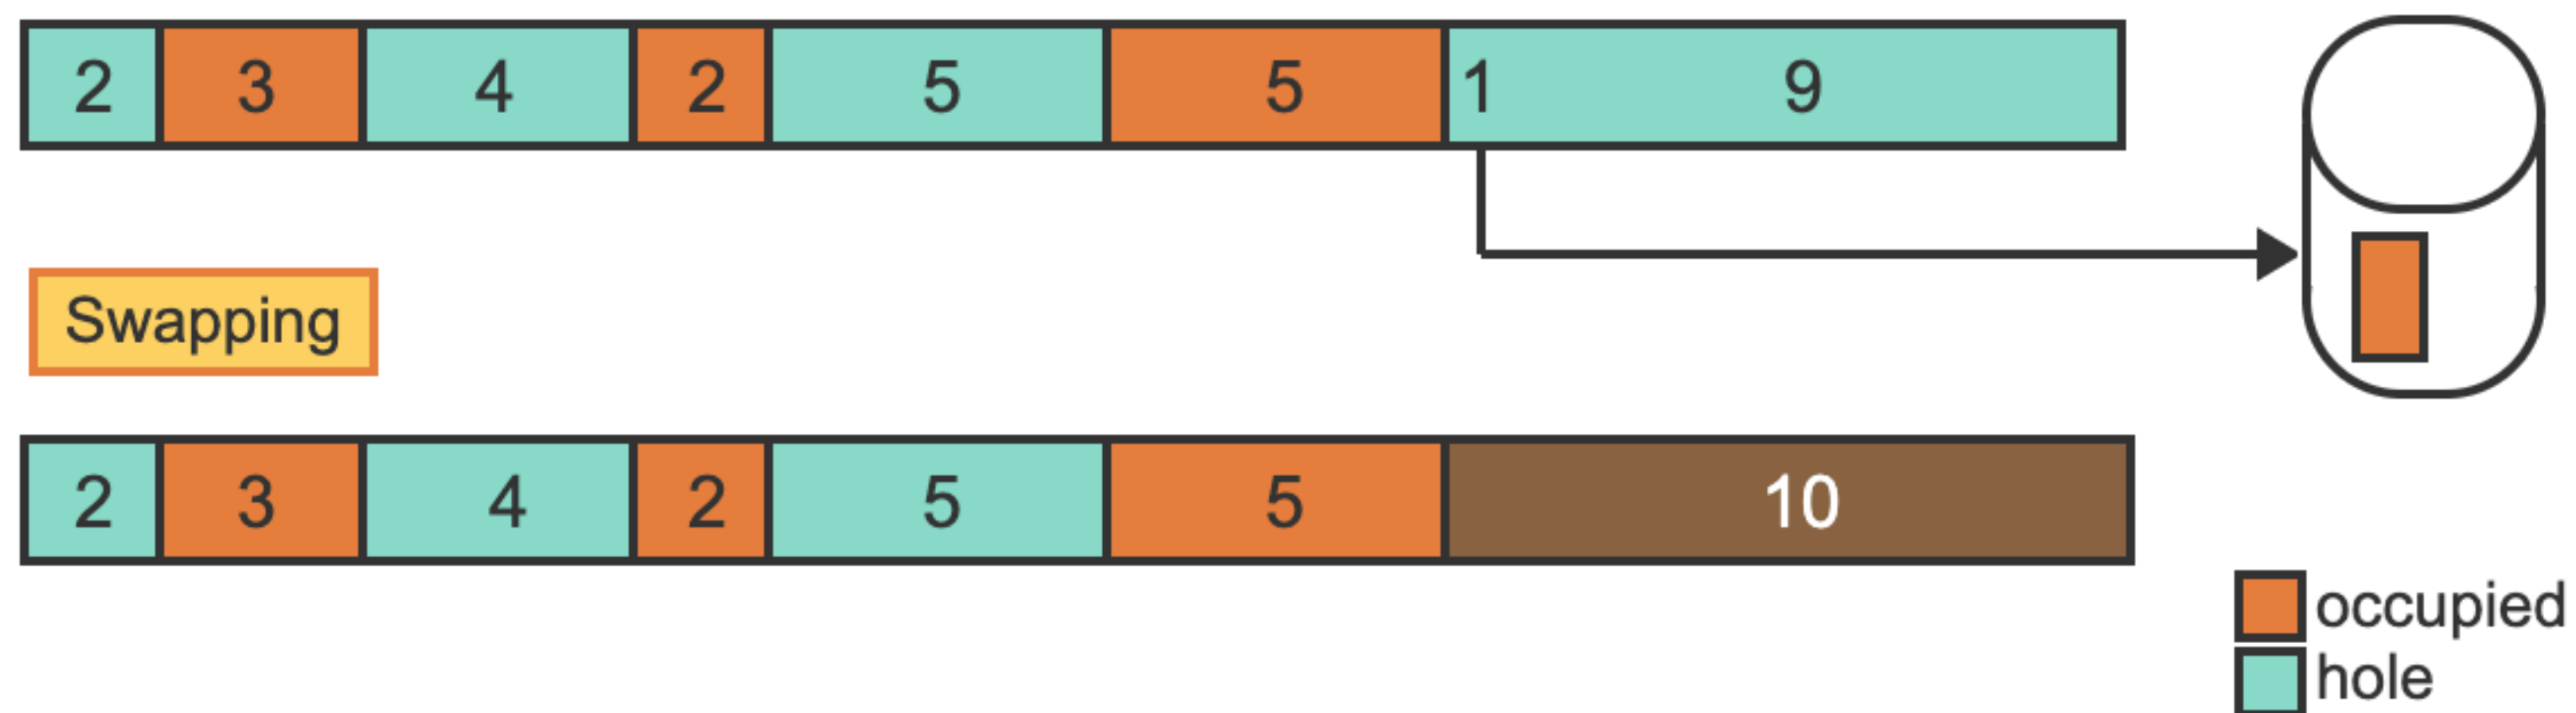
- Current State:



- Current State:



- After Swapping:




- Current State:



• Current State:



• After Compaction:



Paging

What is paging?

- Break physical memory into fixed size blocks: frames
- Break logical memory into same size blocks: pages

Advantages of Paging

- Avoids External Fragmentation
- No need for compaction